Estimating heterogeneous treatment effects versus building individualized treatment rules: Connection and disconnection

Zhongyuan Chen, Jun Xie *

Department of Statistics, Purdue University, 150 N. University Street, West Lafayette, IN 47907, United States of America

ARTICLE INFO

ABSTRACT

Estimating heterogeneous treatment effects is a well studied topic in the statistics literature. More recently, it has regained attention due to an increasing need for precision medicine as well as the increased use of state-of-art machine learning methods in the estimation. Furthermore, estimating heterogeneous treatment effects is directly related to building an individualized treatment rule, which is a decision rule of treatment according to patient characteristics. This paper examines the connection and disconnection between these two research problems. Notably, a better estimation of the heterogeneous treatment effects may or may not lead to a better individualized treatment rule. We provide theoretical frameworks to explain the connection and disconnection and demonstrate two different scenarios through simulations. Our conclusion sheds light on a practical guide that under certain circumstances, there is no need to enhance estimation of the treatment effects, as it does not alter the treatment decision.

© 2023 Elsevier B.V. All rights reserved.

1. Introduction

Estimation of heterogeneous treatment effects (HTE) is a research problem commonly raised in many fields, including politics, economics, education, and healthcare. Instead of an overall treatment effect, heterogeneity exists for subgroups or individuals within a population. We represent HTE by the conditional average treatment effect (CATE) function given a set of covariates. There is a rich literature on estimating the CATE, mostly through regression methods, as well as more recently developed machine learning algorithms (Kunzel et al., 2019). In theory, the accuracy of a CATE estimator is measured by the expected mean squared error (EMSE), which is based on the quadratic loss function. The convergence rate of the EMSE can be used to compare different CATE estimators.

Estimating heterogeneous treatment effects is critical for making medical decisions, such as what treatment to recommend. Building personalized treatment, also referred to as an individualized treatment rule (ITR) (Qian and Murphy, 2011), is usually done through estimating the CATE first and then defining the optimal ITR as the sign function of the CATE (for two treatment options denoted as 1 and −1). Chen et al. (2022) provided a good introduction on building ITRs and called the aforementioned method an indirect approach. Although related, estimating the CATE and building optimal ITRs are separately studied in the literature. In this paper, we elaborate upon the connection and disconnection between the two research areas with mathematical frameworks.

In Section 2, we review recent developments on the estimation of the CATE, including the use of a machine learning method named X-learner (Kunzel et al., 2019), which can improve estimation accuracy with a faster rate of convergence than other conventional estimators. In Section 3, we connect estimating the CATE to building an ITR and show that the

---

* Corresponding author.
  E-mail address: junxie@purdue.edu (J. Xie).

optimal ITR is indeed the sign function of the true CATE function. With this relationship, it is natural to expect that better estimators, in terms of smaller EMSE, lead to improved ITRs. This, however, does not always happen due to the mismatch between the loss functions of these two research problems. More specifically, we use the quadratic loss function for estimation of the CATE but the 0–1 loss function for comparing ITRs. In Section 4, we examine the disconnection through a mathematical framework. That is, for many cases, improving the CATE estimators does not change the corresponding ITRs. Our conclusion sheds light on a practical guide that in certain situations, there is no need to enhance estimation of the CATE, as it does not alter the treatment decision. Section 5 provides simulation examples to display the connection and disconnection. Some discussions are given in Section 6. A proof is given in Appendix.

2. Recent development on estimation of HTE

Suppose we have data from a two-arm clinical trial with (X, A, Y), where $Y \in \mathbb{R}$ denotes a treatment response variable (the larger value the better), $X \in \mathbb{X} \subset \mathbb{R}^p$ is a set of covariates, and $A \in \mathcal{A} = \{-1, 1\}$ denotes the treatment index corresponding to the control or treatment arm. We assume $(X, A, Y) \sim \mathcal{P}$, where $\mathcal{P}$ is a distribution from a specific family. Denote the conditional mean $E(Y|X, A)$. Define

$$\mu_0(x) = E(Y|X = x, A = -1) \quad \text{and} \quad \mu_1(x) = E(Y|X = x, A = 1).$$

Then, the CATE function is $\tau(x) = \mu_1(x) - \mu_0(x)$. Let $\hat{\tau}(x)$ be an estimator of the CATE from a set of independent random data from $\mathcal{P}$. We are interested in estimators with a small expected mean squared error (EMSE):

$$\text{EMSE}(\mathcal{P}, \hat{\tau}) = E[(\hat{\tau}(\mathcal{X}) - \tau(\mathcal{X}))^2],$$

where the expectation is taken over $\hat{\tau}$ and $\mathcal{X}$, which are assumed independent of each other, e.g., $\hat{\tau}$ is estimated from a training data set and $\mathcal{X}$ denotes a new data.

Various methods are available to estimate the CATE. The most commonly used one is to fit regression models for $\hat{\mu}_1(x)$ and $\hat{\mu}_0(x)$, and then $\hat{\tau}(x) = \hat{\mu}_1(x) - \hat{\mu}_0(x)$. As an addition to the rich literature, Chen et al. (2022) considered very general regression models and applied dimension reduction for high-dimensional covariates. Moreover, supervised learning algorithms are used to estimate $\hat{\mu}_1$ and $\hat{\mu}_0$ by the machine learning community (Hu et al., 2021), including Bayesian Additive Regression Trees (BART) (Chipman et al., 2010) and Random Forest (RF) (Wager and Athey, 2018). One of the most recent developments is an algorithm called X-learner (Kunzel et al., 2019), which specifically exploits structural properties of the CATE function for an improved estimator.

Following the results of Kunzel et al. (2019), we use the convergence rates of the EMSE to compare different CATE estimators. Suppose we observe independent and identically distributed data $(X_i, A_i, Y_i) \sim \mathcal{P}, i = 1, \ldots, N$, with m control units and n treated units, $N = m + n$. Let n denote the smaller sample size of the two treatment arms and assume m and n have a similar scale. Most of the estimation methods have a convergence rate that depends on the estimators of $\hat{\mu}_1$ and $\hat{\mu}_0$. For instance, for a parametric distribution family, e.g., μ_0 and μ_1 are parametric functions and X and Y follow parametric distributions, the ordinary least squares estimators of $\hat{\mu}_1$ and $\hat{\mu}_0$ achieve the optimal EMSE minimax rates of $\mathcal{O}(n^{-1})$. Therefore, the corresponding $\hat{\tau} = \hat{\mu}_1 - \hat{\mu}_0$ also has the EMSE rate of $\mathcal{O}(n^{-1})$. In general, for data from a specific distribution family, $\mathcal{P} \in S$, we denote the EMSE rate of $\hat{\tau}$ as $\mathcal{O}(n^{-a_\mu})$, where $0 < a_\mu \leq 1$ is the EMSE rate of $\hat{\mu}_1$ and $\hat{\mu}_0$.

Through imputation of individual treatment effects, Kunzel et al. (2019) developed a new estimation method of the CATE function, namely the X-learner. The X-learner uses the observed data to impute the unobserved individual treatment effects and then directly estimates the CATE function. Denote the X-learner's convergence rate as $\mathcal{O}(n^{-a_\tau})$. When the CATE function has a simpler structure than μ_0 and μ_1, e.g., zero or approximately linear, or the number of observations in one treatment group (usually the control group) is much larger than that in the other, it is proven that $a_\tau > a_\mu$. To conclude, estimators of the CATE can often be improved in terms of a faster convergence rate of the EMSE. There is a clear tendency to choose the estimator with the smaller EMSE.

3. Connection between the two research areas

We have represented HTE by the CATE function, $\tau(x) = E(Y|X = x, A = 1) - E(Y|X = x, A = -1)$. We now introduce ITR, which is a function $d(x) : \mathbb{X} \to \mathcal{A}$. That is, an ITR is a map from the space of covariates, e.g., prognostic variables, to the space of treatments. An optimal ITR, denoted as $d_0(x)$, is the function that gives the highest mean response.

We use the ITR framework from Qian and Murphy (2011), Chen et al. (2022). Recall $(X, A, Y) \sim \mathcal{P}$. For a given ITR d, let $\mathcal{P}^d$ denote the distribution of (X, A, Y) given that $d(X)$ is used to assign treatments. Then, the expected treatment response under the ITR d is

$$\int Y d\mathcal{P}^d = \int Y \frac{d\mathcal{P}^d}{d\mathcal{P}} d\mathcal{P} = \int Y \frac{\mathbb{1}_{d(X)=A}}{p(A|X)} d\mathcal{P} = E\left[Y \frac{\mathbb{1}_{d(X)=A}}{p(A|X)}\right], \tag{1}$$

where $p(A|X)$ is the treatment assignment probability and is assumed > 0. This expectation is called the value function and denoted by $V(d)$. Formally, the optimal ITR d_0 is the rule that maximizes $V(d)$, that is, $d_0 \in \text{argmax}_d V(d)$.

We can obtain the optimal ITR d_0 by estimating the CATE. To see this, we denote $Q_0(X, A) \triangleq E(Y|X, A)$ and $Q_0(x, a)$ is sometimes referred to as the treatment response function. The CATE function is then

$$\tau(x) = Q_0(x, 1) - Q_0(x, -1).$$

Lemma 1. *Assume $p(a|x) > 0$ for $a = 1, -1$ and $x \in \mathbb{X}$. The optimal ITR d_0 is the sign function $d_0(x) = \mathrm{sign}(\tau(x))$, for x such that $\tau(x) \neq 0$.*

Proof. Note that upper case letters are used to denote random variables and lower case letters for values of the random variables. Using Formula (1), we have

$$V(d) = E\left[\frac{\mathbb{1}_{d(X)=A}}{p(A|X)} E[Y|X, A]\right] = E\left[\sum_{a \in \mathcal{A}} \mathbb{1}_{d(X)=a} Q_0(X, a)\right] = E\left[Q_0(X, d(X))\right].$$

The value for the optimal ITR $V(d_0) = E[Q_0(X, d_0(X))] \leq E[\max_{a \in \mathcal{A}} Q_0(X, a)]$. Meanwhile by the definition of d_0,

$$V(d_0) \geq V(d)|_{d(X) \in \mathrm{argmax}_{a \in \mathcal{A}} Q_0(X,a)} = E[\max_{a \in \mathcal{A}} Q_0(X, a)].$$

Thus, $V(d_0) = E[\max_{a \in \mathcal{A}} Q_0(X, a)]$ and the optimal ITR satisfies $d_0(X) = \mathrm{argmax}_{a \in \mathcal{A}} Q_0(X, a)$. In other words, we have

$$d_0(x) = \left\{ \begin{array}{ll} 1, & \text{if } Q_0(x, 1) > Q_0(x, -1), \\ -1, & \text{if } Q_0(x, 1) < Q_0(x, -1). \end{array} \right.$$

That is,

$$d_0(x) = \mathrm{sign}(\tau(x)) \tag{2}$$

for x that $\tau(x) \neq 0$. $\quad \square$

Lemma 1 indicates that building an optimal ITR can be achieved by accurate estimation of the CATE. We can further justify this by providing a relationship between the value function and the estimation error. More specifically, Chen et al. (2022) have showed that for any ITR d, the reduction in value is upper bounded by the estimation error (See Lemma 1 in Chen et al., 2022):

$$V(d_0) - V(d) \leq C \left(E[(Q(X, A) - Q_0(X, A))^2] \right)^{1/2}, \tag{3}$$

where $Q(X, A)$ is a function $Q : \mathbb{X} \times \mathcal{A} \to \mathbb{R}$ such that $d(X) = \mathrm{argmax}_{a \in \mathcal{A}} Q(X, a)$. If we consider $Q(x, a)$ as an estimator of $Q_0(x, a)$, then $\hat{\tau}(x) \triangleq Q(x, 1) - Q(x, -1)$ is the estimated CATE function and we also have the corresponding ITR $d(x) = \mathrm{sign}(\hat{\tau}(x))$.

Intuitively, this upper bound implies that if the EMSE of Q is small, then the corresponding estimated ITR d is close to the optimal ITR d_0 in terms of the value function. Formula (3) along with Lemma 1 explain the connection between estimating the CATE and building an ITR. They support the approach of minimizing the estimation error of the CATE function and then setting the ITR as the sign function of the estimator, i.e., $d(x) = \mathrm{sign}(\hat{\tau}(x))$.

4. Disconnection

Despite the relationship, it is interesting to notice that in Formula (3), minimizing the upper bound is quite different from minimizing the original value function difference, $V(d_0) - V(d)$, as the upper bound may not be very tight. In fact, we will formally demonstrate that improving the estimation of the CATE function (with smaller EMSE) does not necessarily improve the ITR.

Besides the value function, we use a misclassification error to evaluate ITRs. Formula (2) shows the optimal ITR d_0 is $\mathrm{sign}(\tau(x))$, where $\tau(x)$ is the true CATE function. For an estimated optimal ITR, denoted as $\hat{d}(x)$, which is also a binary decision rule, we define the expected misclassification error rate as follows:

$$R(\hat{d}) = P(\hat{d}(X) \neq d_0(X)), \tag{4}$$

where P is the marginal distribution of $X \in \mathbb{X}$ when $(X, A, Y) \sim \mathcal{P}$. This error is based on the 0–1 loss, while the EMSE is based on the quadratic loss. This mismatch is displayed in Fig. 1. Lin (2004) used similar plots to show examples of different loss functions in classification. For simplicity, imagine we have a true positive treatment effect located at 1 on the horizontal axis. Any estimators of the treatment effects located on the positive side of the horizontal axis give correct treatment decisions. That is, the sign functions agree with the sign of the true treatment effect. Analogously, any estimators on the negative side of the axis give incorrect treatment decisions based on the sign functions. Consider two estimators of the treatment effects and the corresponding ITRs from the sign functions. The expected misclassification error rates of the two ITRs will not change if they fall on the same side around 0, even though the estimators give very different quadratic losses.

Given the data $(X_i, A_i, Y_i) \sim \mathcal{P}$, $i = 1, \ldots, N$, with m control units and n treated units, $N = m + n$, we consider two estimators of the CATE, $\hat{\tau}_1$ and $\hat{\tau}_2$. For instance, $\hat{\tau}_1$ denotes the X-learner and $\hat{\tau}_2 = \hat{\mu}_1 - \hat{\mu}_0$ is from the ordinary least squares estimators $\hat{\mu}_1$ and $\hat{\mu}_0$ as discussed in Section 2. There are two corresponding ITRs as the sign functions, denoted as $\hat{d}_1$ and $\hat{d}_2$, respectively. We now evaluate the CATE estimators and the ITRs by comparing them in terms of the EMSE and the expected misclassification error rate, respectively.
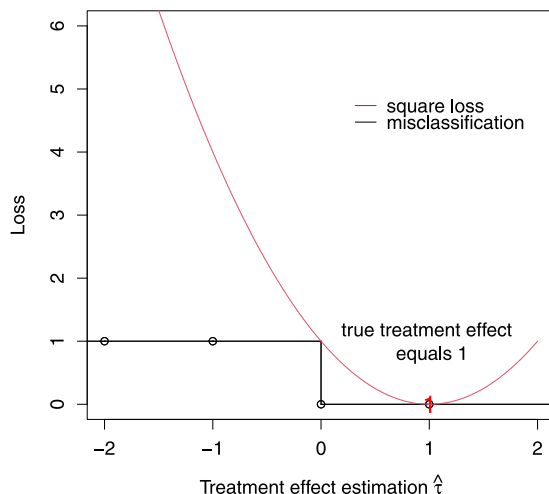
Fig. 1. Mismatch of the two loss functions: square loss and 0–1 loss.

Essentially, for any given estimation method, its performance depends on the specific data distribution $(X_i, A_i, Y_i) \sim \mathcal{P}$. We can talk about a property of $\mathcal{P}$ that will be held for a family of the distributions $\mathcal{P}$. More specifically, given two estimation methods $\hat{\tau}_1$ and $\hat{\tau}_2$ as denoted above, we can define a family of distributions as follows,

$$S = \{\mathcal{P} \text{ such that } \text{EMSE}(\mathcal{P}, \hat{\tau}_1) < \text{EMSE}(\mathcal{P}, \hat{\tau}_2)\}.$$

For example, for distribution families that satisfy Conditions 1–6 as stated in Kunzel et al. (2019), we know the EMSE of the X-learner is smaller than the other standard learner methods as long as the sample sizes m and n are large enough (Theorem 1–2 in Kunzel et al., 2019). In other words, the relationship, $\text{EMSE}(\mathcal{P}, \hat{\tau}_1) < \text{EMSE}(\mathcal{P}, \hat{\tau}_2)$, is valid for a whole distribution family, $\mathcal{P} \in S$, where $\mathcal{P}$ is the underlying data-generating distribution. We further define different distribution families that satisfy different relationships between $\hat{\tau}_1$ and $\hat{\tau}_2$.

Definition 1. Consider $(X_i, A_i, Y_i) \sim \mathcal{P}$, $i = 1, \ldots, N$, with m control units and n treated units and $N = m + n$, $0 < m$, $n < N$. For two estimation strategies of the CATE function, $\hat{\tau}_1$ and $\hat{\tau}_2$, we define three distribution families satisfying the following conditions:

(a) Let S_0 denote the set of all distributions $\mathcal{P}$ such that $\text{EMSE}(\mathcal{P}, \hat{\tau}_1) < \text{EMSE}(\mathcal{P}, \hat{\tau}_2)$.
(b) Let S_1 denote the set of all distributions $\mathcal{P}$ such that $(\hat{\tau}_1(x) - \tau(x))^2 < (\hat{\tau}_2(x) - \tau(x))^2$ for every $x \in \mathbb{X}$. That is, the pointwise squared error of $\hat{\tau}_1$ is smaller than $\hat{\tau}_2$.
(c) Let S_2 denote the set of all distributions $\mathcal{P}$ such that $R(\hat{d}_1) < R(\hat{d}_2)$. That is, the expected misclassification error rate of $\hat{d}_1$ is smaller than $\hat{d}_2$, where $\hat{d}_1(x) = \text{sign}(\hat{\tau}_1(x))$ and $\hat{d}_2(x) = \text{sign}(\hat{\tau}_2(x))$.

Lemma 2. *Consider two estimation strategies for the CATE function, e.g., the X-learner and the standard ordinary least squares estimator denoted as $\hat{\tau}_1$ and $\hat{\tau}_2$. Following Definition 1, we have*

$$S_2 \subset S_1 \subset S_0,$$

where $\subset$ denotes a strict subset.

The proof is provided in Appendix.

Remark 1. Lemma 2 indicates that there are many situations where improving a CATE estimator does not change ITRs. In fact, when $\hat{\tau}_1(x)$ and $\hat{\tau}_2(x)$ have the same sign for $x \in \mathbb{X}$, the corresponding ITRs are exactly the same with no improvement.

Remark 2. Lemma 2 suggests a practice guideline when we apply different estimation methods for the CATE function. We really need to examine the sign functions of different estimators. When there is no change of signs between two estimators, or the changes are minimal, we should choose the estimation method based on the computational cost instead of the EMSE accuracy.

Note that the X-learner is always computationally more expensive than other estimation algorithms, because the X-learner requires several additional computation steps of data imputation. We should not use the X-learner when there is minimal difference between the sign functions of the X-learner and that of the other estimation methods.

Table 1
Two scenario settings in the simulation studies.

Scenario 1			
treatment $A = 1$		control $A = -1$	
range of x	response function	range of x	response function
$-1 < x < -0.4$	$\mu_1(x) = 1.1$	$-1 < x < -0.4$	$\mu_0(x) = 1.0$
$-0.4 < x < 0.5$	$\mu_1(x) = 1.6$	$-0.4 < x < 0.5$	$\mu_0(x) = 1.5$
$0.5 < x < 1.0$	$\mu_1(x) = 1.1$	$0.5 < x < 1.0$	$\mu_0(x) = 1.0$
Scenario 2			
treatment $A = 1$		control $A = -1$	
range of x	response function	range of x	response function
$-1 < x < -0.4$	$\mu_1(x) = 1.6$	$-1 < x < -0.4$	$\mu_0(x) = 1.0$
$-0.4 < x < 0.5$	$\mu_1(x) = 2.1$	$-0.4 < x < 0.5$	$\mu_0(x) = 1.5$
$0.5 < x < 1.0$	$\mu_1(x) = 1.6$	$0.5 < x < 1.0$	$\mu_0(x) = 1.0$

Remark 3. From Definition 1(b), we also have $|\hat{\tau}_1(x) - \tau(x)| < |\hat{\tau}_2(x) - \tau(x)|$ for every $x \in \mathbb{X}$. That is, the mean absolute error of $\hat{\tau}_1$ is smaller than $\hat{\tau}_2$.

5. Simulations

We study two simulation scenarios as specified in Table 1. The first scenario demonstrates a case when a better CATE estimation leads to a better ITR. The second scenario, on the other hand, demonstrates another case when a better estimation does not lead to a better ITR. We generate samples for the treatment and control arms independently. We simulate a covariate $X_i \sim \text{Unif}[-1, 1]$ and then the response $Y_i = \mu(X_i) + \epsilon_i$, where $\mu(x) = \mu_1(x)$ for the treatment arm, and $\mu(x) = \mu_0(x)$ for the control arm, and ϵ_i is the random error following a normal distribution with mean zero and standard deviation 0.01. In the training data, we vary sample sizes of the control group and the treatment group, $m = 200, 400, 600, 800$ and $n = 10, 20, 30, 40$. In the testing data, $m = 1000$ and $n = 50$. Our simulations mimic the situations where the number of samples in the control group is much larger than that in the treatment group.

We consider two machine learning methods for estimating the CATE: the first estimator is the X-learner, denoted as $\hat{\tau}_X$. The second estimator, denoted as $\hat{\tau}$, is a standard estimator $\hat{\tau} = \hat{\mu}_1 - \hat{\mu}_0$. For both estimators, we use the simple linear regression to estimate the response function μ_1 in the treatment group, as the sample size is small. We, however, use a locally weighted regression method, LOESS (Cleveland and S.J., 1988), to estimate μ_0. We compare the two methods across four aspects: (1) MAE: mean absolute error calculated as $\sum_{i=1}^{N} |\hat{\tau}(X_i) - \tau(X_i)|/N$; (2) EMSE: empirical mean squared error calculated as $\sum_{i=1}^{N} (\hat{\tau}(X_i) - \tau(X_i))^2/N$; (3) misclassification error rate: the percentage of the number of misclassified treatments, $\sum_{i=1}^{N} \mathbb{1}(\hat{d}(X_i) \neq d_0(X_i))/N$, where d_0 is the true optimal ITR and $\hat{d}$ is the estimated ITR; (4) empirical value function: an estimate of the value function as in Qian and Murphy (2011), Zhao et al. (2012); and (5) computational time.

The plots on the left column in Fig. 2 (A,C,E,G,I) demonstrate the connection between treatment effect estimation and ITR. In this scenario, a better treatment effect estimation results in a better ITR. In particular, the performances of $\hat{\tau}_X$ are better than $\hat{\tau}$ in terms of smaller MAE, smaller EMSE, lower misclassification rate, and higher value function. The improvements are bigger with larger sample sizes. On the other hand, the plots on the right column in Fig. 2 (B,D,F,H,J) illustrate the disconnection between treatment effect estimation and ITR. That is, although the MAE and EMSE of $\hat{\tau}_X$ are smaller than those of $\hat{\tau}$, the misclassification error rate and the value function are almost the same using $\hat{\tau}_X$ and $\hat{\tau}$. In other words, a better treatment effect estimation does not result in a better ITR in this scenario.

It is important to notice that the computational time of $\hat{\tau}_X$ is longer than that of $\hat{\tau}$ for both scenarios. In the second scenario, since $\hat{\tau}_X$ does not improve ITR, we should not use $\hat{\tau}_X$ but prefer the standard estimation method for ITR, with more efficient computation.

6. Discussion

The literature on estimation of heterogeneous treatment effects mainly uses the EMSE for evaluations of estimation methods. If our ultimate goal is to make treatment recommendations, we show that the evaluation results from the HTE literature are not adequately useful. ITRs are typically dichotomous decision rules. ITRs are assessed by the expected misclassification error, as in Formula (4), or the value function as in Formula (1). Another metric called regret is also used in the literature (Athey and Wager, 2018), which is a modified version of the value function, or a weighted misclassification error. The disconnection between estimating HTEs and building ITRs may be surprising but will enable us to properly choose an estimation method for obtaining the optimal ITR.
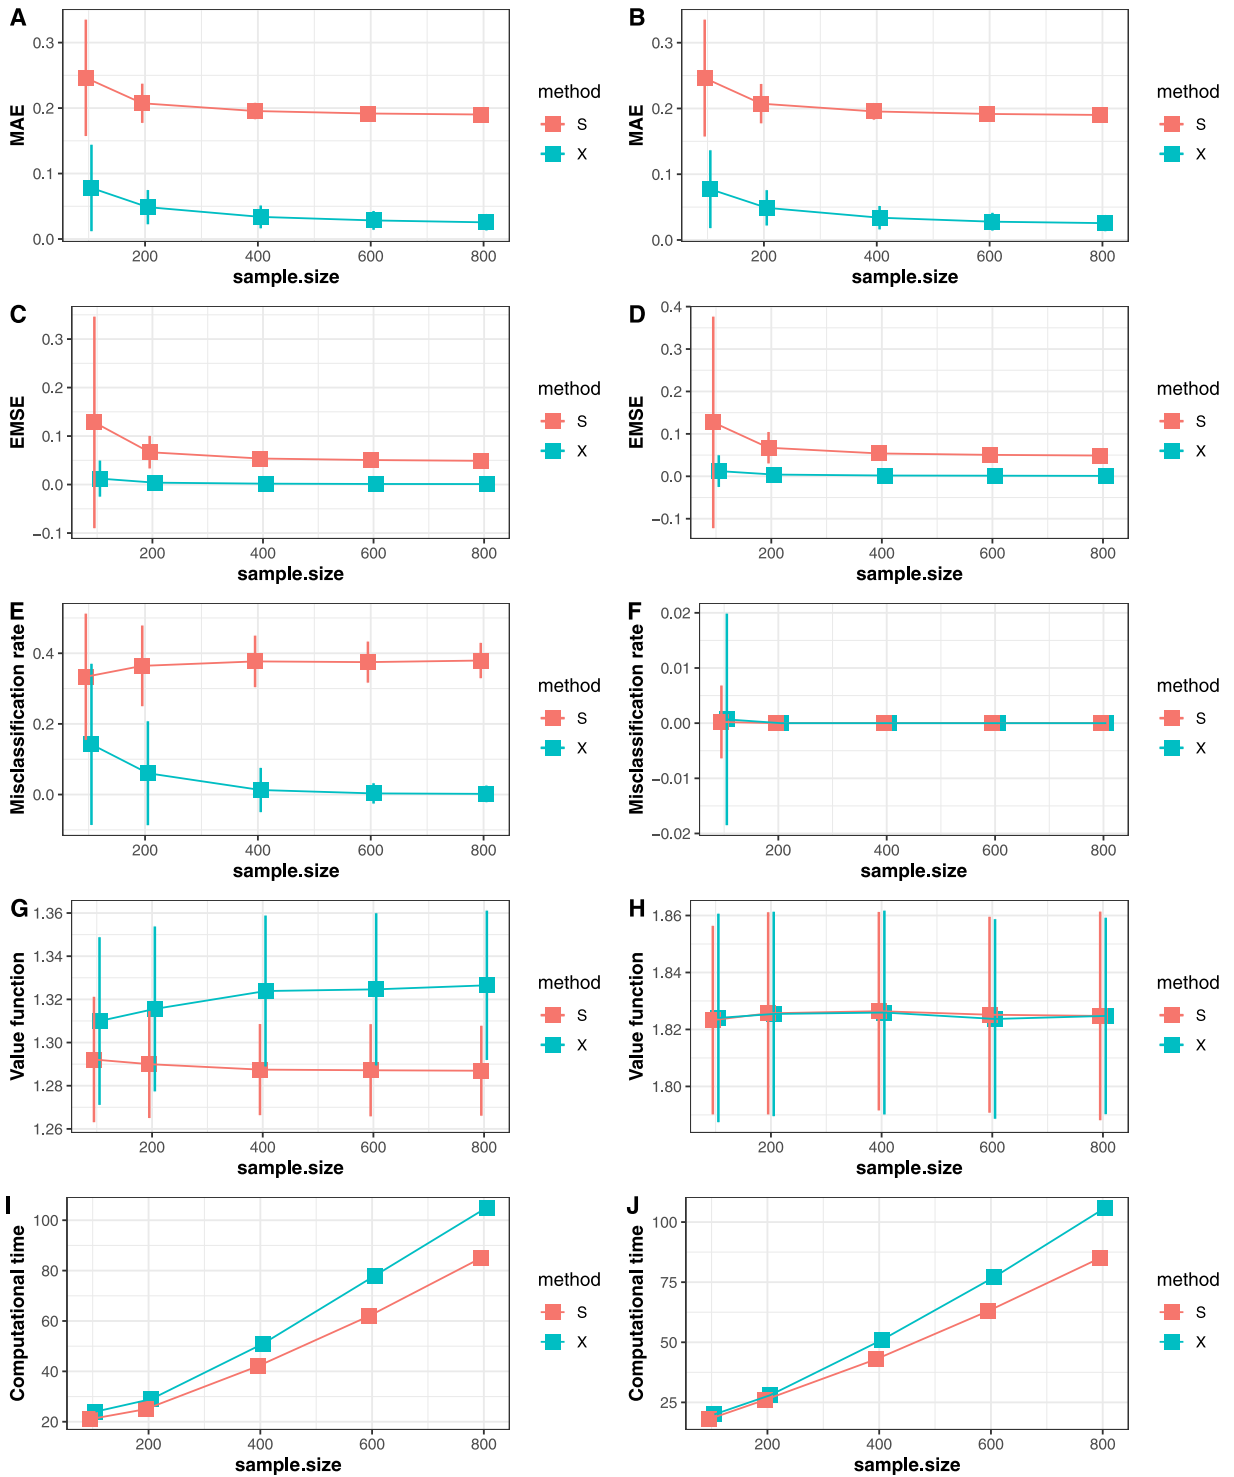
Fig. 2. Comparison of the estimator $\hat{\tau}_X$ (X) from the X-learner and the standard estimator $\hat{\tau}$ (S) in terms of: A-B. Mean absolute error (MAE); C–D. Empirical mean square error (EMSE); E–F. Misclassification error rates; G–H. Empirical value function; I–J. Computational time (unit: seconds) from 1000 simulations. Left column: Scenario 1 (a better HTE estimation leads to a better ITR); right column: Scenario 2 (a better HTE estimation does not lead to a better ITR).

Data availability

No data was used for the research described in the article.

Appendix. Proof of Lemma 2

Proof. Following Definition 1(a) and 1(b), and the EMSE definition, $\text{EMSE}(\mathcal{P}, \hat{\tau}) = E[(\hat{\tau}(\mathcal{X}) - \tau(\mathcal{X}))^2]$, it is straightforward that for any $\mathcal{P} \in S_1$ we must have $\mathcal{P} \in S_0$. Therefore, we obtain $S_1 \subset S_0$ and S_1 as a strict subset of S_0.

Next, for an ITR denoted as $d(x)$ and any $x \in \mathbb{X}$, define

$$R(d(x)) = \begin{cases} 1, & \text{if } d(x) \neq d_0(x), \\ 0, & \text{if } d(x) = d_0(x), \end{cases}$$

where d_0 is the optimal ITR and $d_0(x) = \text{sign}(\tau(x))$. That is, $R(d(x))$ is the pointwise misclassification error of the ITR d at $x \in \mathbb{X}$.

For $\mathcal{P} \in S_1$, from Definition 1(b), we know $(\hat{\tau}_1(x) - \tau(x))^2 < (\hat{\tau}_2(x) - \tau(x))^2$ for any $x \in \mathbb{X}$. Without loss of generality, assume $\tau(x) > 0$. There are three different cases for the misclassification error of $\hat{d}_1(x)$ and $\hat{d}_2(x)$:

1. $R(\hat{d}_1(x)) = R(\hat{d}_2(x))$, if $\hat{\tau}_1(x)$ and $\hat{\tau}_2(x)$ are on the same side of the origin, either both positive or both negative, and $|\hat{\tau}_1(x) - \tau(x)| < |\hat{\tau}_2(x) - \tau(x)|$.
2. $R(\hat{d}_1(x)) > R(\hat{d}_2(x))$, if $|\hat{\tau}_1(x) - \tau(x)| < |\hat{\tau}_2(x) - \tau(x)|$ but $\hat{\tau}_1(x) < 0 < \tau(x) < \hat{\tau}_2(x)$.
3. $R(\hat{d}_1(x)) < R(\hat{d}_2(x))$, if $\hat{\tau}_2(x) < 0 < \hat{\tau}_1(x) < \tau(x)$, or $\hat{\tau}_2(x) < 0 < \tau(x) < \hat{\tau}_1(x)$ and $|\hat{\tau}_1(x) - \tau(x)| < |\hat{\tau}_2(x) - \tau(x)|$.

Integrating $R(\hat{d}_1(x))$ and $R(\hat{d}_2(x))$ over $x \in \mathbb{X}$ with respect of the marginal distribution of X, we obtain the misclassification errors for the two ITRs $\hat{d}_1$ and $\hat{d}_2$, respectively. The distribution family $S_2 = \{\mathcal{P} \text{ such that } R(\hat{d}_1) < R(\hat{d}_2)\}$ corresponds to Item 3 from the list. Therefore, for any $\mathcal{P} \in S_2$ we must have $\mathcal{P} \in S_1$. We obtain $S_2 \subset S_1$ and S_2 as a strict subset. $\square$

References

Athey, S., Wager, S., 2018. Efficient policy learning. arXiv no. 1702.02896.

Chen, Z., Wang, Z., Song, Q., Xie, J., 2022. Data-guided treatment recommendation with feature scores. Statist. Sinica 32, 2497–2519.

Chipman, H.A., George, E.I., Mcculloch, R.E., 2010. BART: Bayesian additive regression trees. Ann. Appl. Stat. 4 (1), 266–298.

Cleveland, W.S., S.J., Devlin, 1988. Locally weighted regression: An approach to regression analysis by local fitting. J. Amer. Statist. Assoc. 83, 596–610.

Hu, L., Ji, J., Li, F., 2021. Estimating heterogeneous survival treatment effect in observational data using machine learning. Stat. Med. 40, 4691–4713.

Kunzel, S.R., Sekhona, J.S., Bickel, P.J., Yu, B., 2019. Metalearners for estimating heterogeneous treatment effects using machine learning. Proc. Natl. Acad. Sci. USA 116, 4156–4165.

Lin, Y., 2004. A note on margin-based loss functions in classification. Statist. Probab. Lett. 68, 73–82.

Qian, M., Murphy, S.A., 2011. Performance guarantees for individualized treatment rules. Ann. Statist. 39, 1180–1210.

Wager, S., Athey, S., 2018. Estimation and inference of heterogeneous treatment effects using random forests. J. Amer. Statist. Assoc. 113, 1228–1242.

Zhao, Y., Zeng, D., Rush, A.J., Kosorok, M.R., 2012. Estimating individualized treatment rules using outcome weighted learning. J. Amer. Statist. Assoc. 107, 1106–1118.